

AMATH 582 Homework 4: Faces and Music Classification

Elise Bishoff

University of Washington

Department of Applied Mathematics

1410 NE Campus Parkway, Seattle WA 98105

Abstract

Singular Value Decomposition poses to be a powerful mathematical tool. In this paper we discuss how we use it to create eigenfaces and reconstruct images of faces based off the different eigenfaces. Also, we use it to perform music classification by taking the SVD of our spectrograms of 5 second clips of songs and then training our classification models on the signatures of the 5 second clips.

Sec. I. Introduction and Overview

Overall, we use Singular Value Decomposition to reconstruct images of faces and classify 5 second clips on song based off their band and music genre. We first discuss the necessary theoretical background needed to understand the problem at hand. We specifically explain Singular Value Decomposition, Support Vector Machines, K-Nearest Neighbors, and Naïve Bayes. Then, we go over the process of applying these methods to the specific problem at hand by exploring the algorithmic development and MATLAB code. This is followed by the computational results for our different scenarios. Next, we have our conclusion. This is followed by Appendices of MATLAB functions used and MATLAB code.

Sec. II. Theoretical Background

Singular Value Decomposition (SVD)

SVD is a very important mathematical tool that allows us to take a matrix and decompose it into three different components that each have specific applications and meaning. The SVD decomposition is illustrated in **Eq 2.1**.

$$A = U\Sigma V^* \quad (2.1)$$

From SVD we get $U \in \mathbb{C}^{m \times m}$ is unitary, $V \in \mathbb{C}^{n \times n}$ is unitary, $\Sigma \in \mathbb{R}^{m \times n}$ is diagonal and $A \in \mathbb{C}^{m \times n}$. We can think of U and V^* as rotations and Σ as a stretch. So, we can think of A as a rotation, stretch(compress), and another rotation.

U is considered the left singular vectors and V^* is considered the right singular vectors. Σ is an ordered diagonal matrix of singular values that tells us what modes of A are most important in reconstructing the matrix A .

There are some important theorems that deal with SVD. We will briefly mention some of them. One of them is that every matrix A has a SVD. No matter what we can perform SVD on any matrix. This is an important point because there are a lot of decompositions that require some criteria to work, but for SVD, it always holds true for every matrix $A \in \mathbb{C}^{m \times n}$. The important differentiator is that we are using two orthogonal basis, U and V^* .

Now, the way to compute the SVD decomposition in practice is to first start with **Eq 2.2** and **Eq 2.3**. Now, since U and V are unitary then we know that $U^*U = I, V^*V = I$. This explains **Eq 2.4**.

$$A^T A = (U\Sigma V^*)^T (U\Sigma V^*) \quad (2.2)$$

$$= V\Sigma U^* U\Sigma V^* \quad (2.3)$$

$$= V\Sigma^2 V^* \quad (2.4)$$

This can be then be written into an eigenvalue problem as shown in **Eq 2.5**.

$$A^T A V = V\Sigma^2 \quad (2.5)$$

Now, with eigenvalue problems we must be careful because we don't always have a solution, but since $A^T A$ is a symmetric matrix, then we are guaranteed to have an eigenvalue decomposition. Thus, we can solve for V in this case. Similarly, we can solve for U by following a similar procedure with AA^T and we get the eigenvalue problem as shown in **Eq 2.6**.

$$AA^T U = U\Sigma^2 \quad (2.6)$$

This is how we extract the matrices U and V . To extract Σ we simply can find the eigenvalues of $A^T A$.

Machine Learning Algorithms: Support Vector Machines, K-Nearest Neighbors, and Naïve Bayes,

Support Vector Machines create hyperplanes that optimize the split between different classes or clusters of data. We can have linear SVM or nonlinear SVM which essentially the difference is between whether we create hyperplanes as our 'lines' to split the data or nonlinear classifiers as 'curves' to split our data. Support vectors are the points of data that touch the edge of our boundaries.

K-Nearest neighbors essentially takes all our training data and uses each labeled point within the model. Then depending on the value of k it classifies the new data based on the new data's k nearest neighbors based off distance. For example, if we use k-nearest neighbor where $k = 1$ then the training data is divided completely. Then, for every new point it finds it nearest neighbor and classifies it as that point. If $k = 5$ then the new data point is determined by a majority vote of the labels of the 5 points.

Naïve Bayes is a simple algorithm that relies on Bayes's theorem. It first computes the prior probability distributions of our training data. Then, when new data is received it estimates the label of the data based off that prior probability distribution.

Sec. III. Algorithm Implementation and Development

Overall, I knew the algorithm needed to be using SVD to find the eigenface for part 1. For part 2 I knew I needed to use SVD to decompose our spectrograms of the 5 second clips of songs and use the signature of the song to create a model that could accurately predict new data's label. I go over the important points of the code and discuss the algorithm implementation and development.

Part 1 Code

The important points from the code here are:

- In **lines 1-50** we are simply loading the data of the cropped and uncropped faces and then save the data as a .mat file to use for later use. In **lines 51-58** I simply take one picture from the cropped faces and one picture from the uncropped faces and plot it in grayscale to get a better understanding of what the pictures look like. In **line 60** I compute the SVD on the cropped faces.
- In **line 63-68** I create the plot of the normalized singular values to see which modes are most important for construction of the image. In **line 70** I create an array of values that tells me what

the percentage of each diagonal value has overall. In **lines 72-86** I plot the first 6 eigenfaces modes to get an understanding of the overall face structure from the pictures. In **lines 88-100** I create a loop to tell me the rank for the cropped images depending on the overall percentage of energy I want.

- In **lines 101-130** I create a plot that reconstructs the face depending on the rank given. In **lines 132-147** I create a plot that compares the original image with the reconstructed image of the face.
- In **lines 150-209** we essentially do the same thing as above, but with uncropped images.

Part 2 Code

I have three different files for the three different cases, but they are essentially the same. I will explain case 1 and then give the differences between case 2 and 3 since case 2 and 3 are almost identical. Also, I created two functions that I explain as well.

From the test case 1 MATLAB code the important points are:

- In **lines 1-43** I read in the different song files. Each song I receive is in stereo mode so I convert it to mono mode by taking the mean between the two arrays after I take the transpose of the song. I calculate how long the song is in seconds as well. In **lines 44-69** I perform the two different functions I wrote on the songs. These functions are explained in greater detail below. Essentially, I split the song and find its spectrogram.
- In **lines 70-77** I combine the data together. Then I perform SVD on the songs. In **lines 81-90** I prep the data by find V' of the training set from the SVD and to project the test songs on the basis of the training set to find its V' . In **lines 91-120** I create an array that hold the true labels of the training and test data.
- In **lines 121-211** I create a KNN, SVM, and a Naïve Bayes model and classify the test data. Then I determine the accuracy of the test data depending on the number of modes used.
- In **lines 212-241** I create various plots to understand the data. I plot the singular values, and the accuracy of the three models together based off the number of modes used.

For the test case 2 and 3 MATLAB code the important points that are different from test case 1 are:

- In **lines 25-27** I chose to only use 400 seconds of each imported set of songs because I downloaded over an hour-long playlist of the artist or genre for each one. This took way too long to process so I extract the first 400 seconds.
- In **lines 53-64** I created the training and test set a little differently. I first put everything together and randomly split the data into 80% training data and 20% testing data.

From the function splitsong and spectransform the important points are:

- Splitsong: In **line 3** the function takes in the time of the song in seconds, the song's amplitude, and the sample rate. It spits out a matrix of 5 second clips of the song(s) pieced together. In **line 4** I create a placeholder for the 5 second clips. Then I determine how many 5 second clips can be extracted from the song in **line 6**. I then have a loop in **lines 12-18** that takes the song and splits it into 5 second clips.
- Spectransform: The function takes in a song or lots of songs in **line 3**. If it is lots of songs, then I loop through all the songs in **lines 5-9** and take the absolute value of their spectrogram in **line 6** and convert it into an vector in **line 7**. Then, I save that vector into a matrix where each column represents a different 5 second clip of a song.

Sec. IV. Computational Results

Part I: Yale Faces B

For the first part we were asked to perform SVD analysis on cropped and uncropped images. First, we are asked to interpret the meaning of U , Σ , and V matrices for our SVD analysis. The U is essentially the basis for our images. This gives us a unique basis in being able to project all our images onto this basis. **Figure 4.1** are the first 6 modes (or basis) of our faces for our cropped and uncropped faces. We can see that they are the skeletons of the averages of all the faces. The cropped faces look like they have more defined facial features whereas the uncropped faces just look like outlines of a head. This is one difference between the cropped and uncropped images.

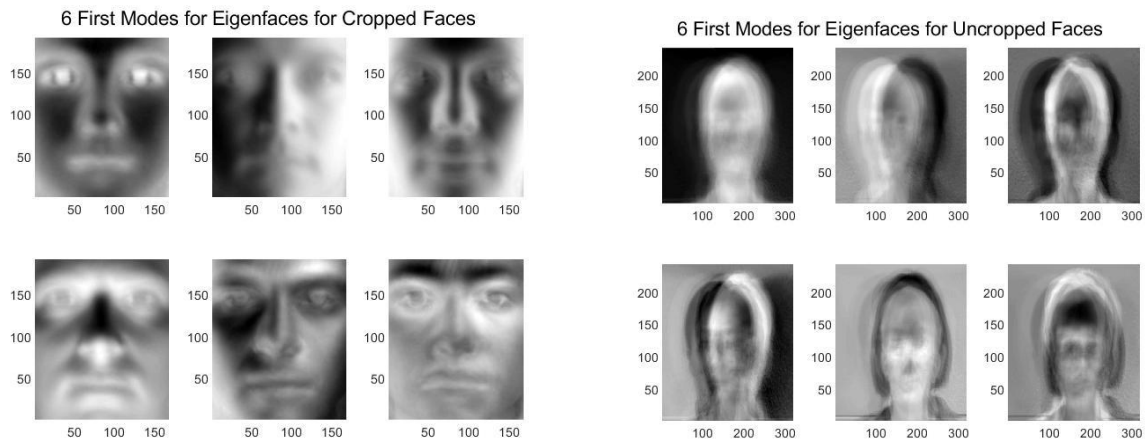


Figure 4.1 The first 6 modes of our faces. Left: Cropped Right: Uncropped

The Σ are the singular values and essentially are ordered from most important to least important. Essentially this tells us which modes are most important for reconstruction of our image. **Figure 4.2** are our Singular Values for our faces. V is the signature of each picture on the basis. It usually is a unique ‘coordinate’ plotted on the basis.

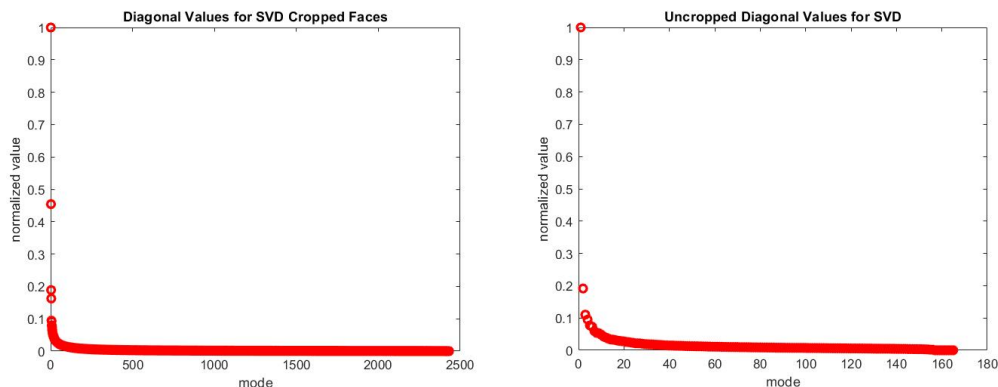


Figure 4.2 The Singular Values for our faces. Left: Cropped Right: Uncropped

Now, in order to get a good image reconstruction, I tested out different rank values. Rank really depends on what you consider to be a good reconstruction. For **Figure 4.3** top image I tried out different rank values depending on how much energy it explained or in other words how many normalized singular

values I needed to add up to get a certain percentage of the total sum of the singular values. **Figure 4.4** shows the number of modes used for each of the images in **Figure 4.3** top in order to get the reconstruction of the image. We can definitely see that the more modes used the better the image.

I decided that rank of 179 which explains 60% of it gives us a good reconstruction of the image which is shown in **Figure 4.3** bottom left. Also, I decided that rank of 94 for uncropped faces worked well to get a reconstruction which is also shown in **Figure 4.3** bottom right. Now, this shows another difference between the cropped and uncropped images. The cropped images I used 179 modes out of around 2400 to explain 60% of it whereas I used 94 modes out around 165 to explain 90% of the image and I still felt like it wasn't as nice of a reconstruction as the uncropped images. This just goes to show that having just the face works a lot better for the SVD than zooming out of the face.

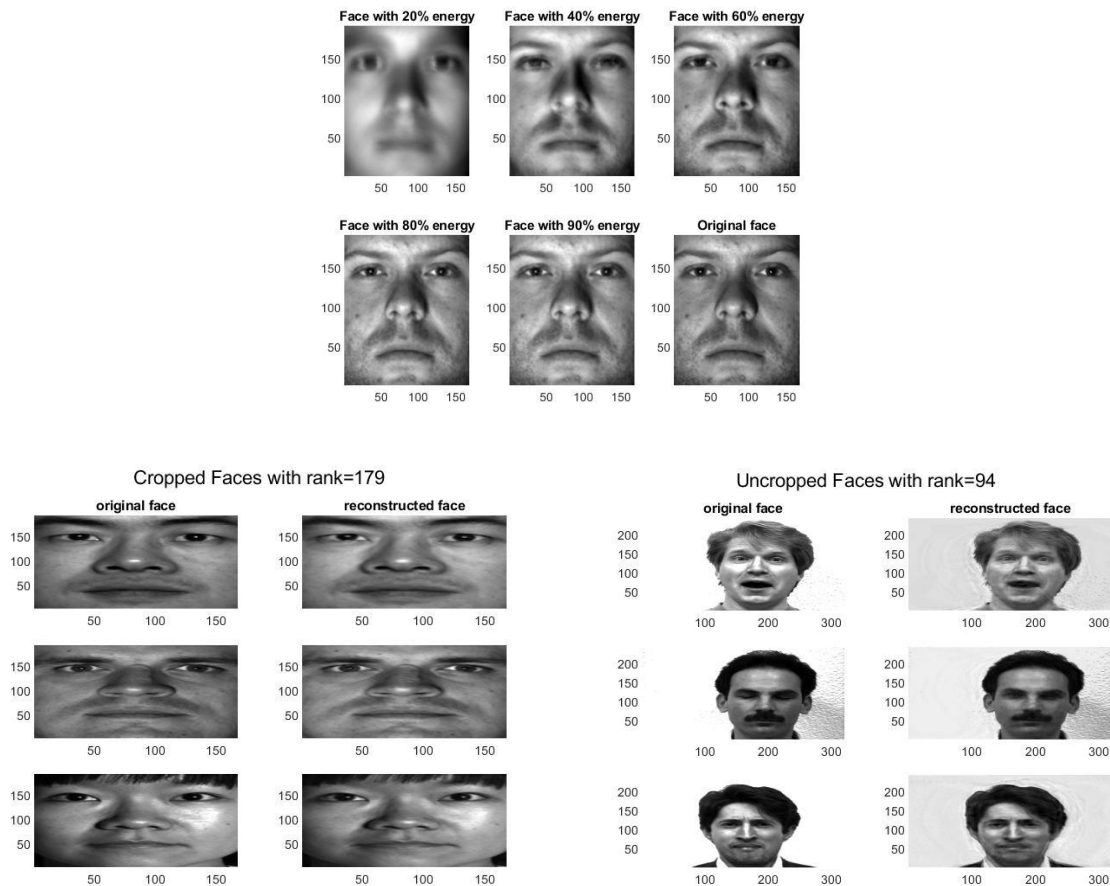


Figure 4.3 Top: Same cropped face reconstructed using different number of modes. Left Bottom: Original image compared to reconstructed cropped faces with 179 modes or rank=179. Right Bottom: Original image compared to reconstructed uncropped faces with rank=94.

Energy	20%	40%	60%	80%	90%
Rank	5	47	179	605	1134

Figure 4.4 Table of the number of modes needed(rank) in order to describe the energy level for Cropped Faces.

Part II: Music Classification

For part 2 we were asked to perform music classification based off three different test cases using SVD. The first case we were asked to perform music classification based off the band. I used Paper Kites, Daft Punk, and Beethoven. The second case we were asked to repeat the same experiment, but with three bands from the same music genre. I used Chopin, Liszt, and Shubert which are all classical composers from roughly the same time period. The third case we were asked to perform music classification based off the genre of music. I used country music, swing music, and classic rock. We use different classification algorithms to train our model and to predict the classes on the test set.

Figure 4.5 shows us the different normalized singular values from each test case. The first test case seems to have less prominent modes meaning that we can gather more information from less modes than the other cases. Whereas case 2 and case 3 seem to have a somewhat less ‘exponential’ curve and tend to be more linear.




Figure 4.5 Singular Values for the different test cases. Left: Test Case 1 Middle: Test Case 2 Right: Test Case 3

Now, I trained with the signature of the songs (V') on three different classification algorithms, K nearest neighbor, Support Vector Machines, and Naïve Bayes. **Figure 4.6** shows the accuracy of classifying our test set based off the different algorithm used and based off how many modes were used in classification.

It was very interesting to see that depending on the type of classification task different classification algorithms worked better than others. For example, Naïve Bayes worked well on test case 1 with more modes but failed on test case 2 and test case 3. Also, SVM did well on test case 1, but flatlined for test case 2 and 3. I copied my code from test case 1 to do test case 2 and 3 so I figured it wasn't the code that was wrong. Then KNN seemed to perform consistently for all test cases. It especially performed well for test case 2 which I thought was odd since it was with three similar bands. All of this can be seen in **Figure 4.6**.

Overall, for test case 1 we were able to get the best accuracy with SVM with around 30 modes with around a 92% accuracy. KNN performed best for test case 2 with getting 100% accuracy with around 55 modes. This was very surprising. Test case 3 performed best with KNN with getting an overall accuracy of around 62% which was surprising as well because I thought it would perform similar to test case 1, but technically it performed the worst out of all the cases.

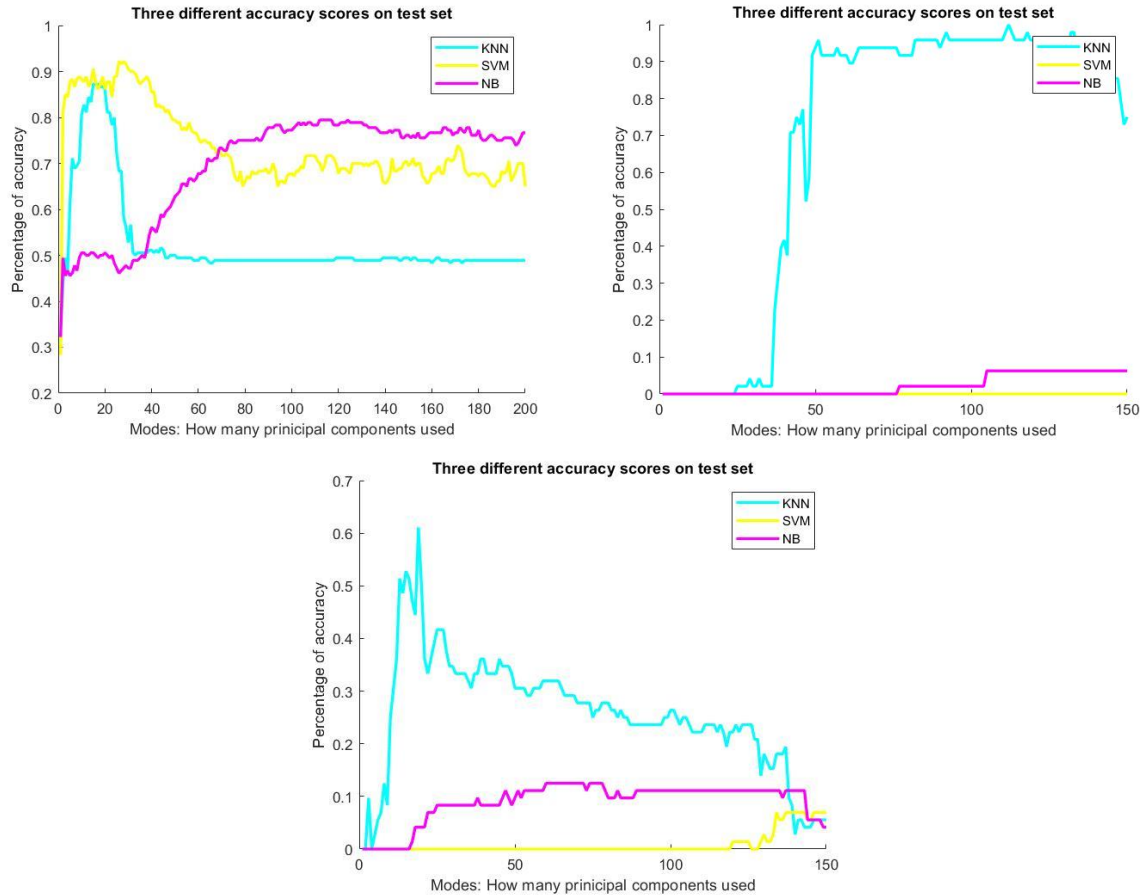


Figure 4.6 The accuracy scores on the test set based off the number of modes used for KNN, SVM, and Naïve Bayes

Sec. V. Summary and Conclusion

In conclusion, we can see that SVD is a powerful tool that allows us to analyze data in new and interesting ways. It allowed for us to decompose and reconstruct images of faces with less information. It allowed for us to construct the ‘average’ face of our dataset. It allowed for us to extract different music signatures and understand that each music signature is unique. It allows for us to classify music based off their band or music genre. These are all things that we were able to accomplish with SVD. Also, we were able to implement different classification algorithms that helped us understand that some algorithms work better than others depending on the problem. Overall, we were once again able to analyze data using powerful tools that allowed us to extrapolate new information from that data that we didn’t see before.

Appendix A MATLAB functions used and brief implementation explanation

audioplayer(v,Fs): This creates an object for *v* where it uses *Fs* as the sample rate. This allowed us to play the audio recording of Handel’s Messiah.

[p,Fs]=audioread(‘musicfile.wav’): This reads data from the file in a format that we can work with in MATLAB.

model = fitcknn(input, input_y): This function allows us to put in our data and its corresponding classes and creates a model using KNN. We can then use this model to predict from other test sets.

model = fitcecoc(input, input_y): This function allows us to put in our data and its corresponding classes and creates a model using SVM. We can then use this model to predict from other test sets.

imshow(image): This displays the image or in our case a frame.

labels = predict(model, test): This function allows us to predict the labels of our test set based off of the given model.

pcolor(0:ss:9,p,spc1.'), shading interp, colormap(hot): pcolor creates a pseudocolor plot. Shading determines how the plot will be colored and colormap defines the color scheme.

y = reshape(s, [l1, l2]): The reshape command allows us to reshape s into a matrix of l1 rows and l2 columns.

s = spectrogram(X): This returns the short-time Fourier transform of the signal X.

[u, s, v] = svd(X): Performs a singular value decomposition on the given matrix, X. s is the diagonal matrix of singular values. Then u and v are either the left singular vectors and right singular vectors.

Appendix B MATLAB code

Part 1

```
1 %For this task we were asked to perform SVD analysis on a set of cropped
2 %and uncropped images.
3 %Here we are just reading in the cropped faces from my computer
4 A=dir('C:\Users\ebish\Documents\MATLAB\582 file 4\CroppedYale\*')
5 croppedfaces=[];
6 for j=3:length(A)
7     d=fullfile(A(j).folder, A(j).name);
8     getpictures=dir(d);
9     for jj=3:length(getpictures)
10         individualpics=fullfile(getpictures(jj).folder,getpictures(jj).name);
11         indpicsfiles=double(imread(individualpics));
12         makevector=indpicsfiles(:);
13         croppedfaces=[croppedfaces makevector];
14     end
15 end
16 %saving the original images dimesions and saving cropped faces so that way
17 %I don't have to run this code everytime and can just load these values
18 %since it takes forever to load.
19 mc=size(indpicsfiles,1);
```



```

20 nc=size(indpicsfiles,2);
21 save('croppedfaces.mat', 'croppedfaces');
22 save('mc.mat', 'mc');
23 save('nc.mat', 'nc');
24 %%
25 %Here we are just reading in the uncropped faces from my computer
26 B=dir('C:\Users\ebish\Documents\MATLAB\582 file 4\yalefaces\*')
27 uncropped=[];
28 for j=3:length(B)
29     dd=fullfile(B(j).folder, B(j).name);
30     unindpicsfiles=double(imread(dd));
31     makevecun=unindpicsfiles(:);
32     uncropped=[uncropped makevecun];
33 end
34 mu=size(unindpicsfiles,1);
35 nu=size(unindpicsfiles,2);
36 %saving the original images dimesions and saving cropped faces so that way
37 %I don't have to run this code everytime and can just load these values
38 %since it takes forever to load.
39 save('uncropped.mat', 'uncropped');
40 save('mu.mat', 'mu');
41 save('nu.mat', 'nu');
42 %%
43 %Load the values from above. Since I ran this code once all I have to do
44 %when i rerun this code is start here. It saves a bunch of time
45 load('uncropped.mat')
46 load('croppedfaces.mat')
47 load('mc')
48 load('nc')
49 load('mu')
50 load('nu')
51 %%
52 %this I am just looking at an image of the cropped and uncropped data
53 t=reshape(croppedfaces(:,10),[mc, nc]);
54 tu=reshape(uncropped(:,10),[mu,nu]);
55 figure
56 pcolor(flip(t)), shading interp, colormap(gray)
57 figure
58 pcolor(flip(tu)), shading interp, colormap(gray)
59 %%
60 %perfomr SVD on cropped images
61 [uc,sigmac,vc]=svd(croppedfaces, 'econ');
62 %%
63 %plot of the modes from most imporant to least important
64 figure
65 plot(diag(sigmac)/max(diag(sigmac)), 'ro', 'Linewidth',[2])
66 xlabel('mode')
67 ylabel('normalized value')
68 title('Diagonal values for SVD Cropped Faces')
69 %This tells us how much each mode is contributing in terms of percentage
70 percentcropped=(diag(sigmac)/max(diag(sigmac)))/sum(diag(sigmac)/max(diag(sigmac)));
71
72 %first 6 eigenfaces
73 figure

```

```

74 subplot(2,3,1)
75 face1=reshape(uc(:,1),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
76 subplot(2,3,2)
77 face1=reshape(uc(:,2),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
78 subplot(2,3,3)
79 face1=reshape(uc(:,3),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
80 subplot(2,3,4)
81 face1=reshape(uc(:,4),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
82 subplot(2,3,5)
83 face1=reshape(uc(:,5),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
84 subplot(2,3,6)
85 face1=reshape(uc(:,6),mc,nc); pcolor(flipud(face1)), shading interp, colormap(gray)
86 sgtitle('6 First Modes for Eigenfaces for Cropped Faces')
87 %%
88 %this allows for me to find the number of modes needed to meet my energy
89 %threshold value of choosing. i chose to do 5 different threshold values
90 modestotal=[];
91 for threshold=[.2 .4 .6 .8 .9]
92     sumperc=0;
93     r=1;
94     while sumperc<threshold
95         sumperc=sumperc+percentcropped(r);
96         r=r+1;
97     end
98     modestotal=[modestotal r];
99 end
100 %%
101 %reconstructing images based off a given number of modes
102 vcprime=vc';
103 modes=modestotal(3);
104 %%
105 %reconstructed images only with different number of modes
106 figure
107 reconstructedimages=uc*sigmac(:,1:modestotal(1))*vcprime(1:modestotal(1),:);
108 subplot(2,3,1)
109 face1=reshape(reconstructedimages(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
110 title('Face with 20% energy')
111 reconstructedimages=uc*sigmac(:,1:modestotal(2))*vcprime(1:modestotal(2),:);
112 subplot(2,3,2)
113 face1=reshape(reconstructedimages(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
114 title('Face with 40% energy')
115 reconstructedimages=uc*sigmac(:,1:modestotal(3))*vcprime(1:modestotal(3),:);
116 subplot(2,3,3)
117 face1=reshape(reconstructedimages(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
118 title('Face with 60% energy')
119 reconstructedimages=uc*sigmac(:,1:modestotal(4))*vcprime(1:modestotal(4),:);
120 subplot(2,3,4)
121 face1=reshape(reconstructedimages(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
122 title('Face with 80% energy')
123 reconstructedimages=uc*sigmac(:,1:modestotal(4))*vcprime(1:modestotal(4),:);

```

```

124 subplot(2,3,5)
125 face1=reshape(reconstructedimages(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
126 title('Face with 90% energy')
127 subplot(2,3,6)
128 face1=reshape(croppedfaces(:,10),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
129 title('Original face')
130 %%
131 %original image compared to reconstructed images
132 figure
133 subplot(3,2,1)
134 face1=reshape(croppedfaces(:,100),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
135 title('original face')
136 subplot(3,2,2)
137 face1=reshape(reconstructedimages(:,100),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
138 title('reconstructed face')
139 subplot(3,2,3)
140 face1=reshape(croppedfaces(:,200),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
141 subplot(3,2,4)
142 face1=reshape(reconstructedimages(:,200),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
143 subplot(3,2,5)
144 face1=reshape(croppedfaces(:,300),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
145 subplot(3,2,6)
146 face1=reshape(reconstructedimages(:,300),mc,nc); pcolor(flipud(face1)), shading interp,
colormap(gray)
147 sgtitle('Cropped Faces with rank=179')
148
149 %% UCropped Section
150 %taking SVD of uncropped images
151 [uu, sigmau, vu]=svd(uncropped, 'econ');
152 %%
153 %plot of the modes from most imporant to least important
154 figure
155 plot(diag(sigmau)/max(diag(sigmau)), 'ro', 'Linewidth',[2])
156 xlabel('mode')
157 ylabel('normalized value')
158 title('Uncropped Diagonal values for SVD')
159 %This tells us how much each mode is contributing in terms of percentage
160 percentuncropped=(diag(sigmau)/max(diag(sigmau)))/sum(diag(sigmau)/max(diag(sigmau)));
161
162 %%
163 %The first 6 eigenfaces
164 figure
165 subplot(2,3,1)
166 face1=reshape(uu(:,1),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
167 subplot(2,3,2)
168 face1=reshape(uu(:,2),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
169 subplot(2,3,3)

```

```

170 face1=reshape(uu(:,3),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
171 subplot(2,3,4)
172 face1=reshape(uu(:,4),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
173 subplot(2,3,5)
174 face1=reshape(uu(:,5),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
175 subplot(2,3,6)
176 face1=reshape(uu(:,6),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
177 sgtitle('6 First Modes for Eigenfaces for Uncropped Faces')
178 %%
179 %this allows for me to find the number of modes needed to meet my energy
180 %threshold value of choosing.
181 sumperc=0;
182 thresholdu=.9;
183 ru=1;
184 while sumperc<thresholdu
185     sumperc=sumperc+percentuncropped(ru);
186     ru=ru+1;
187 end
188 %%reconstructed images only
189 vuprime=vu';
190 modes=ru;
191 reconstructedimagesu=uu*sigmau(:,1:modes)*vuprime(1:modes,:);
192 %%
193 %original compared to reconstructed
194 figure
195 subplot(3,2,1)
196 face1=reshape(uncropped(:,10),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
197 title('original face')
198 subplot(3,2,2)
199 face1=reshape(reconstructedimagesu(:,10),mu,nu); pcolor(flipud(face1)), shading interp,
colormap(gray)
200 title('reconstructed face')
201 subplot(3,2,3)
202 face1=reshape(uncropped(:,20),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
203 subplot(3,2,4)
204 face1=reshape(reconstructedimagesu(:,20),mu,nu); pcolor(flipud(face1)), shading interp,
colormap(gray)
205 subplot(3,2,5)
206 face1=reshape(uncropped(:,30),mu,nu); pcolor(flipud(face1)), shading interp, colormap(gray)
207 subplot(3,2,6)
208 face1=reshape(reconstructedimagesu(:,30),mu,nu); pcolor(flipud(face1)), shading interp,
colormap(gray)
209 sgtitle('Uncropped Faces with rank=94')

```

Part 2 Test Case 1

```

1 %for this assignment we were asked to create a classifier to correctly
2 %classify music files to their artist and genre from 5 second clips.
3 %Here I read in the music, create the files from the spectrogram of the
4 %songs and then use various classifiers to test the different effects on
5 %the different classifiers used
6 %Here we are just reading in the music files.

```

```

7 [p1,Fs1]=audioread('pkbloom2.wav'); %read the audiorecording
8 p1=p1'; %take the transpose
9 p1=mean(p1);
10 tr=length(p1)/Fs1;
11 [p2,Fs2]=audioread('pkfeatherstone.wav');
12 p2=p2';
13 p2=mean(p2);
14 tr2=length(p2)/Fs2;
15 [p3, Fs3]=audioread('dploseyourselftodance.wav');
16 p3=p3';
17 p3=mean(p3);
18 tr3=length(p3)/Fs3;
19 [p4,Fs4]=audioread('dpinstantcrush.wav');
20 p4=p4';
21 p4=mean(p4);
22 tr4=length(p4)/Fs4;
23 [p5, Fs5]=audioread('Beethovenfurelise.wav');
24 p5=p5';
25 p5=mean(p5);
26 tr5=length(p5)/Fs5;
27 [p6,Fs6]=audioread('Beethoven5thSymphony.wav');
28 p6=p6';
29 p6=mean(p6);
30 tr6=length(p6)/Fs6;
31 %reading in the test music files
32 [test1,test1_fs]=audioread('pkwoodland.wav');
33 test1=test1';
34 test1=mean(test1);
35 test1_tr=length(test1)/test1_fs;
36 [test2,test2_fs]=audioread('dpgetlucky.wav');
37 test2=test2';
38 test2=mean(test2);
39 test2_tr=length(test2)/test2_fs;
40 [test3, test3_fs]=audioread('MoonlightSonata.wav');
41 test3=test3';
42 test3=mean(test3);
43 test3_tr=length(test3)/test3_fs;
44 %%
45 %created a function that takes in the song and splits it into 5 seconds
46 %clips. See splitsong function to get better understanding of what it does
47 song1=splitsong(tr,p1,Fs1);
48 song2=splitsong(tr2,p2,Fs2);
49 song3=splitsong(tr3,p3,Fs3);
50 song4=splitsong(tr4,p4,Fs4);
51 song5=splitsong(tr5,p5,Fs5);
52 song6=splitsong(tr6,p6,Fs6);
53 %splitting the songs for the test songs
54 testsong1=splitsong(test1_tr,test1,test1_fs);
55 testsong2=splitsong(test2_tr,test2,test2_fs);
56 testsong3=splitsong(test3_tr,test3,test3_fs);
57 %finding the spectrograms of each song and coverting it to a vector. To find
58 %more information on this function look at spectransform function.
59 specsong1=spectransform(song1);
60 specsong2=spectransform(song2);

```

```

61 specsong3=spectransform(song3);
62 specsong4=spectransform(song4);
63 specsong5=spectransform(song5);
64 specsong6=spectransform(song6);
65 %finding spectrograms for test songs
66 test1_spec=spectransform(testsong1);
67 test2_spec=spectransform(testsong2);
68 test3_spec=spectransform(testsong3);
69 %%
70 %combining the songs into one matrix
71 allsongs=[specsong1 specsong2 specsong3 specsong4 specsong5 specsong6];
72 testsongs=[test1_spec test2_spec test3_spec];
73 %%
74 %finding the mean across rows and then computing the SVD
75 allsongs=allsongs-mean(allsongs,2);
76 [u, sigma, v]=svd(allsongs, 'econ');
77 %%
78 %finding the signature of the songs from the training dataset
79 %and then finding the signature of the test songs on the same basis
80 %as the training dataset.
81 vprime_train=v';
82 vprime_test=inv(sigma)*u'*testsongs;
83 %%
84 %projection of test songs onto the basis for the training songs
85 %found this as well because I wanted to see if it had
86 %any significant impact on the accuracy on the different models
87 Y=u'*testsongs;
88 input=u'*allsongs;
89 Y=Y';
90 input=input';
91 %here we create an array for the actual classes for each 5 second clip and
92 %for loop that simply labels the song with the correct band
93 actual_values=string(zeros(1, size(testsongs,2)));
94 predicted_values=zeros(1, size(testsongs,2));
95 for j=1:length(actual_values)
96     if j<size(testsong1,1)+1
97         actual_values(j)='paperkites';
98     elseif j<size(testsong1,1)+size(testsong2,1)+1
99         actual_values(j)='daftpunk';
100     else
101         actual_values(j)='beethoven';
102     end
103 end
104 %%
105 %this does the same thing as above. IT creates an array with the actual
106 %bands that play each five second clip.
107 band1=size(song1,1)+size(song2,1);
108 band2=size(song3,1)+size(song4,1);
109 band3=size(song5,1)+size(song6,1);
110 actvals_train=string(zeros(1,size(allsongs,2)));
111 for j=1:length(actvals_train)
112     if j<band1+1
113         actvals_train(j)='paperkites';
114     elseif j<band1+band2+1

```

```

115     actvals_train(j)='daftpunk';
116     else
117         actvals_train(j)='beethoven';
118     end
119 end
120 %%
121 %KNN: Uses k nearest neighbor to classify the data. First transformed the
122 %data into rows(where each row is the informaiton on a single 5 second
123 %clip). Then, ran the model over the first 200 modes and created an array
124 %that holds the number of correcly label data from the test set based off
125 %the modes. Then create a plot. A form of validation on the data.
126 vprime_train_normal=vprime_train';
127 vprime_test_normal=vprime_test';
128 actualstrain_prime=actvals_train';
129 corrarray=[];
130 for modes=1:200
131     model=fitcknn(vprime_train_normal(:,1:modes), actvals_train','NumNeighbors',1);
132     label=predict(model,vprime_test_normal(:,1:modes));
133     correctknn=0;
134     for j=1:length(label)
135         if label(j)==actual_values(j)
136             correctknn=correctknn+1;
137         end
138     end
139     corrarray=[corrarray correctknn];
140 end
141 knnperc=corrarray./length(label);
142 figure
143 plot(knnperc,'Linewidth',[2])
144 title('KNN:Accuracy on Test data based off of different modes')
145 xlabel('Modes: How many principal components used')
146 ylabel('Percentage of accuracy')
147 %%
148 %tried the same thing as above out, but used u'data as the data used
149 %instead of just the v from the SVD
150 corrarray2=[];
151 for modes=1:200
152     model=fitcknn(input(:,1:modes), actvals_train','NumNeighbors',1);
153     label=predict(model,Y(:,1:modes));
154     correctknn=0;
155     for j=1:length(label)
156         if label(j)==actual_values(j)
157             correctknn=correctknn+1;
158         end
159     end
160     corrarray2=[corrarray2 correctknn];
161 end
162 figure
163 plot(corrarray2,'Linewidth',[2])
164 %% SMV
165 %SVM: Uses support vector machines to classify the data.Ran the model over
166 %the first 200 modes and created an array
167 %that holds the number of correcly label data from the test set based off
168 %the modes. Then create a plot. A form of validation on the data.

```

```

169 corrarraysvm=[];
170 for modes=1:200
171     modelsvm=fitcecoc(vprime_train_normal(:,1:modes), actvals_train');
172     labelsvm=predict(modelsvm,vprime_test_normal(:,1:modes));
173     correctsvm=0;
174     for j=1:length(labelsvm)
175         if labelsvm(j)==actual_values(j)
176             correctsvm=correctsvm+1;
177         end
178     end
179     corrarraysvm=[corrarraysvm correctsvm];
180 end
181 %percentage of accuracy based off the modes
182 svmperc=corrarraysvm./length(labelsvm);
183 figure
184 plot(svmperc,'Linewidth',[2])
185 title('SVM:Accuracy on Test data based off of different modes')
186 xlabel('Modes: How many principal components used')
187 ylabel('Percentage of accuracy')
188 %% Naïve Bayes
189 %NB: Uses Naïve Bayes to classify the data.Ran the model over
190 %the first 200 modes and created an array
191 %that holds the number of correctly label data from the test set based off
192 %the modes. Then create a plot. A form of validation on the data.
193 corrarraynb=[];
194 for modes=1:200
195     modelnb=fitcnb(vprime_train_normal(:,1:modes), actvals_train');
196     labelnb=predict(modelnb,vprime_test_normal(:,1:modes));
197     correctnb=0;
198     for j=1:length(labelnb)
199         if labelnb(j)==actual_values(j)
200             correctnb=correctnb+1;
201         end
202     end
203     corrarraynb=[corrarraynb correctnb];
204 end
205 %percentage of accuracy based off the modes
206 tperc=corrarraynb./length(labelnb);
207 figure
208 plot(tperc,'Linewidth',[2])
209 title('NB:Accuracy on Test data based off of different modes')
210 xlabel('Modes: How many principal components used')
211 ylabel('Percentage of accuracy')
212 %%
213 %plot of the actual data normalized
214 figure
215 plot(allsongs(:,200), 'Linewidth',[2])
216 xlabel('time')
217 ylabel('y-value')
218 title('The different points extracted at each frame')
219 %%
220 %this looks at the diagonal values. This tells
221 %me what is going on in what direction. If it is big
222 %a lot is happening and if it is small then not

```



```

223 %much is happening.
224 figure
225 plot(diag(sigma)/max(diag(sigma)), 'ro', 'Linewidth',[2])
226 xlabel('mode')
227 ylabel('normalized value')
228 title('Diagonal values for SVD for Test Case 1')
229 %This tells us how much each mode is contributing in terms of percentage
230 percentofenergy=(diag(sigma)/max(diag(sigma)))/sum(diag(sigma)/max(diag(sigma)));
231 %% Plot all plots together of the results
232 figure
233 hold on
234 plot(knnperc,'c', 'Linewidth',[2])
235 plot(svmperc,'y', 'Linewidth',[2])
236 plot(tperc,'m','Linewidth',[2])
237 hold off
238 legend('KNN','SVM','NB')
239 title('Three different accuracy scores on test set')
240 xlabel('Modes: How many principal components used')
241 ylabel('Percentage of accuracy')

```

Part 2 Test Case 2

```

1 %for this assignment we were asked to create a classifier to correctly
2 %classify music files to their artist and genre from 5 second clips.
3 %Here I read in the music, create the files from the spectrogram of the
4 %songs and then use various classifiers to test the different effects on
5 %the different classifiers used
6 %Here we are just reading in the music files.
7 [p1,Fs1]=audioread('schubert.wav'); %read the audiorecording
8 p1=p1'; %take the transpose
9 p1=mean(p1);
10 tr=length(p1)/Fs1;
11 [p2,Fs2]=audioread('liszt.wav');
12 p2=p2';
13 p2=mean(p2);
14 tr2=length(p2)/Fs2;
15 [p3, Fs3]=audioread('chopin.wav');
16 p3=p3';
17 p3=mean(p3);
18 tr3=length(p3)/Fs3;
19 %%
20 %created a function that takes in the song and splits it into 5 seconds
21 %clips. See splitsong function to get better understanding of what it does
22 %Instead of using tr I hard code a value because the actual files is hours
23 %long and takes wayyyyyyy tooooo long to run so I truncate the files to
24 %only give me a smaller chunk of time.
25 song1=splitsong(400,p1,Fs1);
26 song2=splitsong(400,p2,Fs2);
27 song3=splitsong(400,p3,Fs3);
28 %%
29 %finding the spectrograms of each song and coverting it to a vector. To find

```

```

30 %more information on this function look at spectransform function.
31 specsong1=spectransform(song1);
32 specsong2=spectransform(song2);
33 specsong3=spectransform(song3);
34 %this does the same thing as above. IT creates an array with the actual
35 %bands that play each five second clip.
36 band1=size(song1,1);
37 band2=size(song2,1);
38 band3=size(song3,1);
39 actvals_train=string(zeros(1,band1+band2+band3));
40 for j=1:length(actvals_train)
41     if j<band1+1
42         actvals_train(j)='schubert';
43     elseif j<band1+band2+1
44         actvals_train(j)='liszt';
45     else
46         actvals_train(j)='chopin';
47     end
48 end
49 %%
50 %combining the songs into one matrix
51 allsongs=[specsong1 specsong2 specsong3];
52 %%
53 %In this section I simply split my data into training and test data
54 everything=[allsongs; actvals_train];
55 ind=randperm(size(everything,2));
56 everythingmixed=everything(:,ind);
57 %%
58 split=fix(.7*size(everything,2));
59 allsongs_almost=everything(:,1:split);
60 testsongs_almost=everything(:,split+1:end);
61 actual_values=testsongs_almost(end,:);
62 actvals_train=allsongs_almost(end,:);
63 allsongs=double(allsongs_almost(1:end-1,:));
64 testsongs=double(testsongs_almost(1:end-1,:));
65 %%
66 %finding the mean across rows and then computing the SVD
67 allsongs=allsongs-mean(allsongs,2);
68 [u, sigma, v]=svd(allsongs, 'econ');
69 %%
70 %here we create an array for the actual classes for each 5 second clip and
71 %for loop that simply labels the song with the correct band
72 % actual_values=string(zeros(1, size(testsongs,2)));
73 % predicted_values=zeros(1, size(testsongs,2));
74 % for j=1:length(actual_values)
75 %     if j<size(testsong1,1)+1
76 %         actual_values(j)='schubert';
77 %     elseif j<size(testsong1,1)+size(testsong2,1)+1
78 %         actual_values(j)='liszt';
79 %     else
80 %         actual_values(j)='chopin';
81 %     end
82 % end
83 %%

```

```

84 %finding the signature of the songs from the training dataset
85 %and then finding the signature of the test songs on the same basis
86 %as the training dataset.
87 vprime_train=v';
88 vprime_test=inv(sigma)*u'*testsongs;
89 %%
90 %projection of test songs onto the basis for the training songs
91 %found this as well becuse I wanted to see if it had
92 %any significant impact on the accuracy on the different models
93 Y=u'*testsongs;
94 input=u'*allsongs;
95 Y=Y';
96 input=input';
97 %%
98 %KNN: Uses k nearest neighbor to classify the data. First transformed the
99 %data into rows(where each row is the informaiton on a single 5 second
100 %clip). Then, ran the model over the first 200 modes and created an array
101 %that holds the number of correcly label data from the test set based off
102 %the modes. Then create a plot. A form of validation on the data.
103 vprime_train_normal=vprime_train';
104 vprime_test_normal=vprime_test';
105 actualstrain_prime=actuals_train';
106 corrray=[];
107 for modes=1:150
108     model=fitcknn(vprime_train_normal(:,1:modes), actuals_train','NumNeighbors',1);
109     label=predict(model,vprime_test_normal(:,1:modes));
110     correctknn=0;
111     for j=1:length(label)
112         if label(j)==actual_values(j)
113             correctknn=correctknn+1;
114         end
115     end
116     corrray=[corrray correctknn];
117 end
118 knnperc=corrray./length(label);
119 figure
120 plot(knnperc,'Linewidth',[2])
121 title('KNN:Accuracy on Test data based off of different modes')
122 xlabel('Modes: How many principal components used')
123 ylabel('Percentage of accuracy')
124 %%
125 %tried the same thing as above out, but used u'data as the data used
126 %instead of just the v from the SVD
127 corrray2=[];
128 for modes=1:150
129     model=fitcknn(input(:,1:modes), actuals_train','NumNeighbors',1);
130     label=predict(model,Y(:,1:modes));
131     correctknn=0;
132     for j=1:length(label)
133         if label(j)==actual_values(j)
134             correctknn=correctknn+1;
135         end
136     end
137     corrray2=[corrray2 correctknn];

```

```

138 end
139 figure
140 plot(corrarray2,'Linewidth',[2])
141 %% SMV
142 %SVM: Uses support vector machines to classify the data.Ran the model over
143 %the first 200 modes and created an array
144 %that holds the number of correctly label data from the test set based off
145 %the modes. Then create a plot. A form of validation on the data.
146 corrarraysvm=[];
147 for modes=1:150
148     modelsvm=fitcecoc(vprime_train_normal(:,1:modes), actvals_train');
149     labelsvm=predict(modelsvm,vprime_test_normal(:,1:modes));
150     correctsvm=0;
151     for j=1:length(labelsvm)
152         if labelsvm(j)==actual_values(j)
153             correctsvm=correctsvm+1;
154         end
155     end
156     corrarraysvm=[corrarraysvm correctsvm];
157 end
158 %percentage of accuracy based off the modes
159 svmperc=corrarraysvm./length(labelsvm);
160 figure
161 plot(svmperc,'Linewidth',[2])
162 title('SVM:Accuracy on Test data based off of different modes')
163 xlabel('Modes: How many principal components used')
164 ylabel('Percentage of accuracy')
165 %% Naive Bayes
166 %NB: Uses Naive Bayes to classify the data.Ran the model over
167 %the first 200 modes and created an array
168 %that holds the number of correctly label data from the test set based off
169 %the modes. Then create a plot. A form of validation on the data.
170 corrarraynb=[];
171 for modes=1:150
172     modelnb=fitcnb(vprime_train_normal(:,1:modes), actvals_train');
173     labelnb=predict(modelnb,vprime_test_normal(:,1:modes));
174     correctnb=0;
175     for j=1:length(labelnb)
176         if labelnb(j)==actual_values(j)
177             correctnb=correctnb+1;
178         end
179     end
180     corrarraynb=[corrarraynb correctnb];
181 end
182 %percentage of accuracy based off the modes
183 tperc=corrarraynb./length(labelnb);
184 figure
185 plot(tperc,'Linewidth',[2])
186 title('NB:Accuracy on Test data based off of different modes')
187 xlabel('Modes: How many principal components used')
188 ylabel('Percentage of accuracy')
189 %%
190 %plot of the actual data normalized
191 figure

```

```

192 plot(allsongs(:,200), 'Linewidth',[2])
193 xlabel('time')
194 ylabel('y-value')
195 title('The different points extracted at each frame')
196 %%
197 %this looks at the diagonal values. This tells
198 %me what is going on in what direction. If it is big
199 %a lot is happening and if it is small then not
200 %much is happening.
201 figure
202 plot(diag(sigma)/max(diag(sigma)), 'ro', 'Linewidth',[2])
203 xlabel('mode')
204 ylabel('normalized value')
205 title('Diagonal values for SVD Case 2')
206 %This tells us how much each mode is contributing in terms of percentage
207 percentofenergy=(diag(sigma)/max(diag(sigma)))/sum(diag(sigma)/max(diag(sigma)));
208 %% Plot all plots together of the results
209 figure
210 hold on
211 plot(knnperc,'c', 'Linewidth',[2])
212 plot(svmperc,'y', 'Linewidth',[2])
213 plot(tperc,'m','Linewidth',[2])
214 hold off
215 legend('KNN','SVM','NB')
216 title('Three different accuracy scores on test set')
217 xlabel('Modes: How many principal components used')
218 ylabel('Percentage of accuracy')
219

```

Part 2 Test Case 3

```

1 %for this assignment we were asked to create a classifier to correctly
2 %classify music files to their artist and genre from 5 second clips.
3 %Here I read in the music, create the files from the spectrogram of the
4 %songs and then use various classifiers to test the different effects on
5 %the different classifiers used
6 %Here we are just reading in the music files.
7 [p1,Fs1]=audioread('rockclassic.wav'); %read the audiorecording
8 p1=p1'; %take the transpose
9 p1=mean(p1);
10 tr=length(p1)/Fs1;
11 [p2,Fs2]=audioread('country.wav');
12 p2=p2';
13 p2=mean(p2);
14 tr2=length(p2)/Fs2;
15 [p3, Fs3]=audioread('swing.wav');
16 p3=p3';
17 p3=mean(p3);
18 tr3=length(p3)/Fs3;
19 %%
20 %created a function that takes in the song and splits it into 5 seconds
21 %clips. See splitsong function to get better understanding of what it does
22 %Instead of using tr I hard code a value because the actual files is hours

```

```

23 %long and takes wayyyyyyy toooooo long to run so I truncate the files to
24 %only give me a smaller chunk of time.
25 song1=splitsong(400,p1,Fs1);
26 song2=splitsong(400,p2,Fs2);
27 song3=splitsong(400,p3,Fs3);
28 %%
29 %finding the spectrograms of each song and coverting it to a vector. To find
30 %more information on this function look at spectransform function.
31 specsong1=spectransform(song1);
32 specsong2=spectransform(song2);
33 specsong3=spectransform(song3);
34 %this does the same thing as above. IT creates an array with the actual
35 %bands that play each five second clip.
36 band1=size(song1,1);
37 band2=size(song2,1);
38 band3=size(song3,1);
39 actvals_train=string(zeros(1,band1+band2+band3));
40 for j=1:length(actvals_train)
41     if j<band1+1
42         actvals_train(j)='rock';
43     elseif j<band1+band2+1
44         actvals_train(j)='liszt';
45     else
46         actvals_train(j)='chopin';
47     end
48 end
49 %%
50 %combining the songs into one matrix
51 allsongs=[specsong1 specsong2 specsong3];
52 %%
53 %In this section I simply split my data into training and test data
54 everything=[allsongs; actvals_train];
55 ind=randperm(size(everything,2));
56 everythingmixed=everything(:,ind);
57 %%
58 split=fix(.8*size(everything,2));
59 allsongs_almost=everything(:,1:split);
60 testsongs_almost=everything(:,split+1:end);
61 actual_values=testsongs_almost(end,:);
62 actvals_train=allsongs_almost(end,:);
63 allsongs=double(allsongs_almost(1:end-1,:));
64 testsongs=double(testsongs_almost(1:end-1,:));
65 %%
66 %finding the mean across rows and then computing the SVD
67 allsongs=allsongs-mean(allsongs,2);
68 [u, sigma, v]=svd(allsongs, 'econ');
69 %%
70 %here we create an array for the actual classes for each 5 second clip and
71 %for loop that simply labels the song with the correct band
72 % actual_values=string(zeros(1, size(testsongs,2)));
73 % predicted_values=zeros(1, size(testsongs,2));
74 % for j=1:length(actual_values)
75 %     if j<size(testsong1,1)+1
76 %         actual_values(j)='schubert';

```

```

77 %     elseif j<size(testsong1,1)+size(testsong2,1)+1
78 %         actual_values(j)='liszt';
79 %     else
80 %         actual_values(j)='chopin';
81 %     end
82 % end
83 %%
84 %finding the signature of the songs from the training dataset
85 %and then finding the signature of the test songs on the same basis
86 %as the training dataset.
87 vprime_train=v';
88 vprime_test=inv(sigma)*u'*testsongs;
89 %%
90 %projection of test songs onto the basis for the training songs
91 %found this as well because I wanted to see if it had
92 %any significant impact on the accuracy on the different models
93 Y=u'*testsongs;
94 input=u'*allsongs;
95 Y=Y';
96 input=input';
97 %%
98 %KNN: Uses k nearest neighbor to classify the data. First transformed the
99 %data into rows(where each row is the information on a single 5 second
100 %clip). Then, ran the model over the first 200 modes and created an array
101 %that holds the number of correctly label data from the test set based off
102 %the modes. Then create a plot. A form of validation on the data.
103 vprime_train_normal=vprime_train';
104 vprime_test_normal=vprime_test';
105 actualstrain_prime=actuals_train';
106 corrray=[];
107 for modes=1:150
108     model=fitcknn(vprime_train_normal(:,1:modes), actuals_train', 'NumNeighbors',1);
109     label=predict(model,vprime_test_normal(:,1:modes));
110     correctknn=0;
111     for j=1:length(label)
112         if label(j)==actual_values(j)
113             correctknn=correctknn+1;
114         end
115     end
116     corrray=[corrray correctknn];
117 end
118 knnperc=corrray./length(label);
119 figure
120 plot(knnperc,'Linewidth',[2])
121 title('KNN:Accuracy on Test data based off of different modes')
122 xlabel('Modes: How many principal components used')
123 ylabel('Percentage of accuracy')
124 %%
125 %tried the same thing as above out, but used u'data as the data used
126 %instead of just the v from the SVD
127 corrray2=[];
128 for modes=1:150
129     model=fitcknn(input(:,1:modes), actuals_train', 'NumNeighbors',1);
130     label=predict(model,Y(:,1:modes));

```

```

131     correctknn=0;
132     for j=1:length(label)
133         if label(j)==actual_values(j)
134             correctknn=correctknn+1;
135         end
136     end
137     corrray2=[corrray2 correctknn];
138 end
139 figure
140 plot(corrray2,'Linewidth',[2])
141 %% SMV
142 %SVM: Uses support vector machines to classify the data.Ran the model over
143 %the first 200 modes and created an array
144 %that holds the number of correctly label data from the test set based off
145 %the modes. Then create a plot. A form of validation on the data.
146 corrraysvm=[];
147 for modes=1:150
148     modelsvm=fitcecoc(vprime_train_normal(:,1:modes), actvals_train');
149     labelsvm=predict(modelsvm,vprime_test_normal(:,1:modes));
150     correctsvm=0;
151     for j=1:length(labelsvm)
152         if labelsvm(j)==actual_values(j)
153             correctsvm=correctsvm+1;
154         end
155     end
156     corrraysvm=[corrraysvm correctsvm];
157 end
158 %percentage of accuracy based off the modes
159 svmperc=corrraysvm./length(labelsvm);
160 figure
161 plot(svmperc,'Linewidth',[2])
162 title('SVM:Accuracy on Test data based off of different modes')
163 xlabel('Modes: How many principal components used')
164 ylabel('Percentage of accuracy')
165 %% Naive Bayes
166 %NB: Uses Naive Bayes to classify the data.Ran the model over
167 %the first 200 modes and created an array
168 %that holds the number of correctly label data from the test set based off
169 %the modes. Then create a plot. A form of validation on the data.
170 corrraynb=[];
171 for modes=1:150
172     modelnb=fitcnb(vprime_train_normal(:,1:modes), actvals_train');
173     labelnb=predict(modelnb,vprime_test_normal(:,1:modes));
174     correctnb=0;
175     for j=1:length(labelnb)
176         if labelnb(j)==actual_values(j)
177             correctnb=correctnb+1;
178         end
179     end
180     corrraynb=[corrraynb correctnb];
181 end
182 %percentage of accuracy based off the modes
183 tperc=corrraynb./length(labelnb);
184 figure

```



```

185 plot(tperc,'Linewidth',[2])
186 title('NB:Accuracy on Test data based off of different modes')
187 xlabel('Modes: How many principal components used')
188 ylabel('Percentage of accuracy')
189 %%
190 %plot of the actual data normalized
191 figure
192 plot(allsongs(:,150), 'Linewidth',[2])
193 xlabel('time')
194 ylabel('y-value')
195 title('The different points extracted at each frame')
196 %%
197 %this looks at the diagonal values. This tells
198 %me what is going on in what direction. If it is big
199 %a lot is happening and if it is small then not
200 %much is happening.
201 figure
202 plot(diag(sigma)/max(diag(sigma)), 'ro', 'Linewidth',[2])
203 xlabel('mode')
204 ylabel('normalized value')
205 title('Diagonal values for SVD for test case 3')
206 %This tells us how much each mode is contributing in terms of percentage
207 percentofenergy=(diag(sigma)/max(diag(sigma)))/sum(diag(sigma)/max(diag(sigma)));
208 %% Plot all plots together of the results
209 figure
210 hold on
211 plot(knnperc,'c', 'Linewidth',[2])
212 plot(svmperc,'y', 'Linewidth',[2])
213 plot(tperc,'m','Linewidth',[2])
214 hold off
215 legend('KNN','SVM','NB')
216 title('Three different accuracy scores on test set')
217 xlabel('Modes: How many principal components used')
218 ylabel('Percentage of accuracy')
219

```

Spectrogram Function

```

1 %this function takes a song and saves the absolute value spectrogram
2 %information in a vector and appends it to an array called songlist
3 function songlist=spectransform(song)
4 songlist=[];
5 for j=1:size(song,1)
6     s=abs(spectrogram(song(j,:)));
7     s=reshape(s,[size(s,1)*size(s,2),1]);
8     songlist=[songlist s];
9 end

```

SplitSong Function

```

1 %this function takes in the time of the song,tr, the song amplitude
2 %file,py, and the sample rate, Fsp and chops the song into 5 second clips

```

```
3 function songs=splitsong(tr,py,Fsp)
4 songs=[]; %file to hold the 5 second clips
5 %this tells me how many 5 second clips I can get from the song
6 totalsongclips=fix(tr/5);
7 %initial start and end for the first song
8 starttemp=1*Fsp;
9 endtemp=6*Fsp;
10 %loop that take a song and splits it into 5 second clips and then holds it
11 %in the array, songs.
12 for j=1:totalsongclips-1 %I take one less just to avoid complications with
13     %how long we have at the end of the song
14     segtemp=py(:,starttemp:endtemp);
15     songs=[songs; segtemp];
16     starttemp=starttemp+5*Fsp;
17     endtemp=endtemp+5*Fsp;
18 end
```